


Llama-2-13B FP16 · SLO attainment on NRP

Homogeneous · 2×A10, one node Heterogeneous · A10 + RTX 5000 Ada, two sites



Each point is 1 of 100 measured requests; all requests succeeded. 32 generated tokens; one sweep per arm. The dotted line marks 99% attainment. Deadlines are absolute seconds because the paper's A100 latency baseline was not measured here. The two 13B layouts differ in GPU memory, placement, and parallelism; this is not a cost-matched 70B replication.